


# Gradient Boosting

## Gradient Boosted Decision Trees



Gradient Boosting is an ensemble learning technique that builds a strong predictive model by combining multiple weak learners in an iterative fashion. It works by sequentially adding new models that focus on correcting the errors of the previous models, gradually improving the overall performance.

# Ensemble Learning and Weak Learners

## Ensemble Learning

Gradient Boosting is a type of ensemble learning, which combines multiple models to achieve better performance than any individual model.


## Weak Learners

Gradient Boosting uses weak learners, such as decision trees, as the building blocks to construct a stronger, more accurate model.

## Iterative Improvement

By adding new weak learners that focus on correcting the errors of the previous models, Gradient Boosting iteratively improves the overall prediction.

# Gradient Boosting Algorithm



1

## Step 1

Start with a basic weak learner, such as a decision tree with a single split.

2

## Step 2


Compute the residuals (errors) between the current model's predictions and the true target values.

3

## Step 3

Train a new weak learner to predict the residuals and add it to the ensemble.

# Minimizing Errors through Iterative Adjustments



1

## Residual Computation

Calculate the errors between the current model's predictions and the true target values.

2

## Model Adjustment

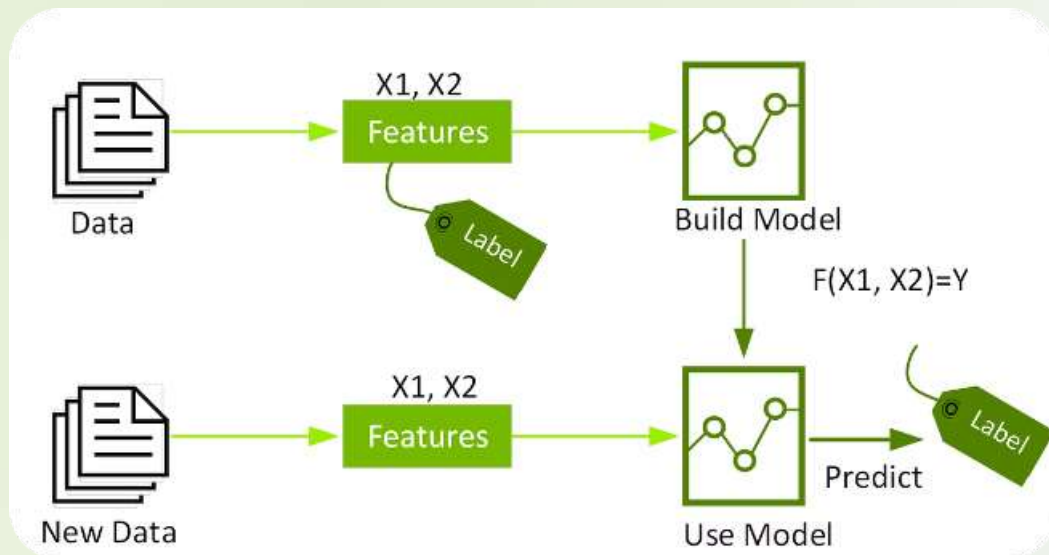
Train a new weak learner to predict the residuals and add it to the ensemble.

3

## Error Reduction

The new model reduces the overall error, improving the ensemble's predictive performance.

# Advantages of Gradient Boosting



## 1 High Accuracy

Gradient Boosting can achieve state-of-the-art performance on a wide range of machine learning problems.

## 2 Robustness


It is resilient to outliers and can handle both numerical and categorical features.

## 3 Interpretability

The individual weak learners in the ensemble can provide insights into the underlying patterns in the data.

## 4 Versatility

Gradient Boosting can be applied to classification, regression, and ranking tasks.



# Hyperparameter Tuning for Gradient Boosting

## Learning Rate

Determines the contribution of each new weak learner to the ensemble. A lower rate can lead to better generalization.

## Maximum Depth

Controls the complexity of the individual weak learners. Deeper trees can capture more complex patterns.

## Number of Estimators

Specifies the number of weak learners to be added to the ensemble. More estimators can improve performance.

## Regularization

Techniques such as L1 or L2 regularization can help prevent overfitting.

# Applications of Gradient Boosting



## Classification

Gradient Boosting is widely used for classification tasks, such as spam detection, sentiment analysis, and credit risk assessment.



## Regression

It is also effective for regression problems, including stock price prediction, sales forecasting, and energy demand forecasting.



## Ranking

Gradient Boosting algorithms can be used for ranking tasks, such as search engine result ranking and recommendation systems.





# Conclusion and Key Takeaways

1

## Ensemble Learning

Gradient Boosting is a powerful ensemble learning technique that combines multiple weak learners to create a strong predictive model.

2

## Iterative Improvement

The algorithm iteratively adjusts the model by training new weak learners to correct the errors of the previous models.

3

## Versatility and Performance

Gradient Boosting can be applied to a wide range of machine learning tasks and has been shown to deliver state-of-the-art performance.

4

## Hyperparameter Tuning

Careful selection of hyperparameters, such as learning rate and maximum depth, is crucial for optimizing the performance of Gradient Boosting models.

