

CODE FOR THE GIVEN PROBLEM STATEMENT

Importing necessary libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split,
GridSearchCV
from sklearn.preprocessing import StandardScaler,
LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report,
confusion_matrix, roc_auc_score, roc_curve
```

Step 1: Load the data

Load the data from the uploaded CSV file

```
data
=pd.read_csv('/mnt/data/customer_purchase_data.csv')
```

Step 2: Explore the dataset

```
print("First few rows of the dataset:")
print(data.head()) # Display first few rows of the dataset
```

```
print("\nDataset information:")
print(data.info()) # Summary of the dataset
```

```
print("\nStatistical summary of the dataset:")
print(data.describe()) # Statistical summary of numerical
columns
```

```
print("\nChecking for missing values:")
```

```
print(data.isnull().sum()) # Checking for missing values
```

```
# Step 3: Data Preprocessing
```

```
# Handling missing values (example: filling missing  
numerical values with mean)
```

```
data.fillna(data.mean(), inplace=True)
```

```
# Encoding categorical variables
```

```
label_encoders = { }
```

```
for column in
```

```
data.select_dtypes(include=['object']).columns:
```

```
    le = LabelEncoder()
```

```
    data[column] = le.fit_transform(data[column])
```

```
    label_encoders[column] = le
```

```
# Replace 'target_column' with the actual target column  
name from your dataset
```

```
PurchaseStatus= 'PurchaseStatus' # Adjust this with the  
actual column name
```

```
X = data.drop(PurchaseStatus, axis=1)
```

```
y = data[PurchaseStatus]
```

```
# Train-test split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

```
# Feature Scaling
```

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
# Step 4: Model Development
```

```
# Using RandomForestClassifier as an example
```

```
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
```

Step 5: Model Evaluation

```
y_pred = model.predict(X_test)
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

```
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

ROC-AUC Score

```
roc_auc = roc_auc_score(y_test,
model.predict_proba(X_test)[:, 1])
print(f"\nROC-AUC Score: {roc_auc}')
```

Plotting the ROC Curve

```
fpr, tpr, thresholds = roc_curve(y_test,
model.predict_proba(X_test)[:, 1])
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'AUC = {roc_auc:.2f}')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.grid()
plt.show()
```

Step 6: Hyperparameter Tuning

Example: Tuning Random Forest hyperparameters

```
param_grid = {
    'n_estimators': [100, 200, 300],
```

```

    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10]
}

```

```

grid_search = GridSearchCV(estimator=model,
param_grid=param_grid, cv=5, scoring='roc_auc',
n_jobs=-1)
grid_search.fit(X_train, y_train)

```

Best Parameters from Grid Search

```

print("\nBest Parameters found by Grid Search:")
print(grid_search.best_params_)

```

Best Model Evaluation

```

best_model = grid_search.best_estimator_
y_pred_best = best_model.predict(X_test)
print("\nBest Model Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_best))

```

```

print("\nBest Model Classification Report:")
print(classification_report(y_test, y_pred_best))

```

Best Model ROC-AUC Score

```

roc_auc_best = roc_auc_score(y_test,
best_model.predict_proba(X_test)[:, 1])
print(f'\nBest Model ROC-AUC Score: {roc_auc_best}')

```

SOLUTION OR OUTPUT FOR THE PROBLEM STATEMENT

First few rows of the dataset:

	Age	Gender	AnnualIncome	NumberOfPurchases	ProductCategory \
0	40	1	66120.267939	8	0

1	20	1	23579.773583	4	2
2	27	1	127821.306432	11	2
3	24	1	137798.623120	19	3
4	31	1	99300.964220	19	1

	TimeSpentOnWebsite	LoyaltyProgram	DiscountsAvailed	PurchaseStatus
0	30.568601	0	5	1
1	38.240097	0	5	0
2	31.633212	1	0	1
3	46.167059	0	4	1
4	19.823592	0	0	1

Dataset information:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1500 entries, 0 to 1499

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Age	1500 non-null	int64
1	Gender	1500 non-null	int64
2	AnnualIncome	1500 non-null	float64
3	NumberOfPurchases	1500 non-null	int64
4	ProductCategory	1500 non-null	int64
5	TimeSpentOnWebsite	1500 non-null	float64
6	LoyaltyProgram	1500 non-null	int64
7	DiscountsAvailed	1500 non-null	int64
8	PurchaseStatus	1500 non-null	int64

dtypes: float64(2), int64(7)
memory usage: 105.6 KB
None

Statistical summary of the dataset:

	Age	Gender	AnnualIncome	NumberOfPurchases \
count	1500.000000	1500.000000	1500.000000	1500.000000
mean	44.298667	0.504667	84249.164338	10.420000
std	15.537259	0.500145	37629.493078	5.887391
min	18.000000	0.000000	20001.512518	0.000000
25%	31.000000	0.000000	53028.979155	5.000000
50%	45.000000	1.000000	83699.581476	11.000000
75%	57.000000	1.000000	117167.772858	15.000000
max	70.000000	1.000000	149785.176481	20.000000

	ProductCategory	TimeSpentOnWebsite	LoyaltyProgram	DiscountsAvailed \
count	1500.000000	1500.000000	1500.000000	1500.000000
mean	2.012667	30.469040	0.326667	2.555333
std	1.428005	16.984392	0.469151	1.705152
min	0.000000	1.037023	0.000000	0.000000
25%	1.000000	16.156700	0.000000	1.000000
50%	2.000000	30.939516	0.000000	3.000000
75%	3.000000	44.369863	1.000000	4.000000
max	4.000000	59.991105	1.000000	5.000000

	PurchaseStatus
count	1500.000000
mean	0.432000
std	0.495520
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

Checking for missing values:

Age	0
Gender	0
AnnualIncome	0
NumberOfPurchases	0
ProductCategory	0
TimeSpentOnWebsite	0
LoyaltyProgram	0
DiscountsAvailed	0
PurchaseStatus	0

dtype: int64

Confusion Matrix:

```
[[171  1]
 [ 13 115]]
```


Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.93	0.99	0.96	172
1	0.99	0.90	0.94	128

accuracy			0.95	300
macro avg	0.96	0.95	0.95	300
weighted avg	0.96	0.95	0.95	300

ROC-AUC Score: 0.9556458938953489



Best Parameters found by Grid Search:

{'max_depth': None, 'min_samples_split': 5, 'n_estimators': 200}

Best Model Confusion Matrix:

```
[[169  3]
 [ 13 115]]
```

Best Model Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.93	0.98	0.95	172
1	0.97	0.90	0.93	128

accuracy		0.95		300
macro avg	0.95	0.94	0.94	300
weighted avg	0.95	0.95	0.95	300

Best Model ROC-AUC Score: 0.9543059593023255